

# RMSC4002 Tutorial 9

## Chapter 5-6

November 21, 2017

## 1 Dummy Variable in Logistic Regression

```
> g<-(d3$ln_MV>9.4766)+1 # create dummy var g=2 if ln_MV>9.4766; g=1 otherwise
>summary(glm(HSI~EY+CFTP+g+DY+BTME+DTE+EY*g+CFTP*g+DY*g+BTME*g+DTE*g,
data=d3, binomial))
```

```
Call:
glm(formula = HSI ~ EY + CFTP + g + DY + BTME + DTE + EY * g +
      CFTP * g + DY * g + BTME * g + DTE * g, family = binomial,
      data = d3)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )	
(Intercept)	-21.0469	5.6643	-3.716	0.000203	***
EY	29.0955	38.2277	0.761	0.446592	
CFTP	-0.1916	3.8931	-0.049	0.960745	
g	15.2686	5.3888	2.833	0.004606	**
DY	1.0979	0.8182	1.342	0.179637	
BTME	1.4157	1.5673	0.903	0.366398	
DTE	-0.3659	0.8731	-0.419	0.675166	
EY:g	-28.8939	38.2187	-0.756	0.449641	
CFTP:g	0.1269	3.6225	0.035	0.972054	
g:DY	-0.9840	0.7821	-1.258	0.208316	
g:BTME	-1.3220	1.4948	-0.884	0.376499	
g:DTE	0.2058	0.5665	0.363	0.716429	

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

```
> summary(glm(HSI~EY+g+DY+BTME+DTE+EY*g+DY*g+BTME*g+DTE*g,data=d3,
binomial))
```

```
Call:
glm(formula = HSI ~ EY + g + DY + BTME + DTE + EY * g + DY *
      g + BTME * g + DTE * g, family = binomial, data = d3)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )	
(Intercept)	-21.0169	5.5409	-3.793	0.000149	***
EY	28.6141	27.2787	1.049	0.294200	
g	15.2425	5.2592	2.898	0.003752	**
DY	1.0968	0.7752	1.415	0.157089	
BTME	1.4089	1.4991	0.940	0.347279	
DTE	-0.3558	0.8474	-0.420	0.674536	
EY:g	-28.4191	27.2658	-1.042	0.297274	
g:DY	-0.9858	0.7396	-1.333	0.182540	
g:BTME	-1.3146	1.4225	-0.924	0.355404	
g:DTE	0.2015	0.5513	0.366	0.714692	

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

```
> summary(glm(HSI~g+DY+DY*g,data=d3,binomial))
```

Call:

```
glm(formula = HSI ~ g + DY + DY * g, family = binomial, data = d3)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )	
(Intercept)	-18.6304	3.4117	-5.461	4.74e-08	***
g	12.9236	3.1251	4.135	3.54e-05	***
DY	1.4036	0.6485	2.165	0.0304	*
g:DY	-1.2860	0.6082	-2.115	0.0345	*

---  
Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

```
> lregl<-glm(HSI~g+DY+DY*g,data=d3,binomial)
```

```
> pr<-(lregl$fit>0.5)
```

```
> table(pr,d3$HSI)
```

pr	0	1
FALSE	624	3
TRUE	2	29

The interpretation of this threshold-type logistic model is as follow:

$$\ln[\pi/(1-\pi)] = -18.6304 + 12.9236g + 1.4036DY - 1.286(g * DY)$$

$$\Leftrightarrow \begin{cases} \ln[\pi/(1-\pi)] = (-18.6304 + 12.936) + (1.4036 - 1.286)DY, & \text{for } g = 1 \\ \quad \quad \quad = -5.6944 + 0.1176DY \\ \ln[\pi/(1-\pi)] = (-18.6304 + 12.9236 \times 2) + (1.4036 - 1.286 \times 2)DY, & \text{for } g = 2 \\ \quad \quad \quad = 7.2416 - 1.1684DY \end{cases}$$

where  $\pi$  = probability of HSI=1. Note that for  $g=2$ , the intercept is much bigger than that of  $g=1$ , hence the probability of HSI=1 is higher.

## 2 Classification Tree

```
> d<-read.csv("fin-ratio.csv") # read in data in csv format
> library(rpart) # load rpart library
> ctree<-rpart(HSI~EY+CFTP+ln_MV+DY+BTME+DTE,data=d,method="class")
> plot(ctree,asp=0.5,main=" Fin-ratio" ) # plot ctree
> text(ctree,use.n=T,cex=0.6) # add text
```

```
> print(ctree)
n= 680
node), split, n, loss, yval, (yprob)
* denotes terminal node
1) root 680 32 0 (0.952941176 0.047058824)
2) d$ln_MV< 9.4776 647 3 0 (0.995363215 0.004636785) *
3) d$ln_MV>=9.4776 33 4 1 (0.121212121 0.878787879) *
```

```
# plot ln_MV versus HSI with color, red=group 0 and blue=group 1
> plot(d$HSI,d$ln_MV,pch=21,bg=c("red","blue")[d$HSI+1])
> abline(h=9.478) # add a horizontal line at y=9.478
```

```


> pr<-predict(ctree)           # pr has 2 columns of prob. in group 0 or 1
> cl<-0*(pr[,1]>0.5)+1*(pr[,2]>0.5) # assign group label if prob>0.5
> table(cl,d$HSI)              # cross tabulation

cl  0  1
   0 644  3
   1   4 29

```

**Remarks:**

- “asp”: numeric, giving the aspect ratio  $y/x$ .
- “text”: labels the current plot of the tree dendrogram with text.
- “use.n”: logical, if true, adds to label (#event level 1/#event level 1/etc.)
- “cex”: number indicating the amount by which plotting text and symbols should be scaled relative to the default. 1 is default; 1.5 is 50% larger; 0.5 is 50% smaller, etc.
- “pch”: specifying symbols to use when plotting points as shown below. For symbols 21 through 25, specify border color (col=) and fill color (bg=).



```

> d<-read.csv("iris.csv")
> library(rpart)
> names(d)
[1] "Sepal_len" "Sepal_wid" "Petal_len" "Petal_wid" "Species"
> ctree<-rpart(Species~Sepal_len+Sepal_wid+Petal_len+Petal_wid,data=d,method="class")
> plot(ctree,asp=2)
> text(ctree,use.n=T,cex=0.6)

```

```

> print(ctree)
n= 150

node), split, n, loss, yval, (yprob)
      * denotes terminal node

1) root 150 100 1 (0.33333333 0.33333333 0.33333333)
2) Petal_len< 2.45 50 0 1 (1.00000000 0.00000000 0.00000000) *
3) Petal_len>=2.45 100 50 2 (0.00000000 0.50000000 0.50000000)
   6) Petal_wid< 1.75 54 5 2 (0.00000000 0.90740741 0.09259259) *
   7) Petal_wid>=1.75 46 1 3 (0.00000000 0.02173913 0.97826087) *

```

```

# plot Petal_wid versus Petal_len with different color for each species
> plot(d$Petal_len,d$Petal_wid,pch=21,bg=c("red","blue","green")[d$Species])
> abline(h=1.75)           # add a horizontal line
> abline(v=2.45)           # add a vertical line

```

```

> pr<-predict(ctree)           # pr has 3 columns of prob.
> cl<-max.col(pr)              # cl is the col. no. such that pr is the maximum.
> table(cl,d$Species)         # cross tabulation table

cl  1  2  3
   1 50  0  0
   2  0 49  5
   3  0  1 45

```

**Remarks:**

- “max.col”: Find the maximum position for each row of a matrix.

### 3 Tree Quality assessment

- Support = number of cases in that rule / total number of data in dataset
- Confidence = number of correct-classified data in the node / number of data at the node
- Capture = number of correct-classified data in the node / number of data of this class in the whole dataset

**Exercise 2014-15 final Q5** Classification tree is built using *rpart()* and the result is

```
> print(ctree)
node), split, n, loss, yval, (yprob)
      * denotes terminal node
1) root 620 # 0 (# #)
  2) Bank< 2.5 # # 0 (# #)
    4) Employ< 1.2075 286 52 0 (# #) *
    5) Employ>=1.2075 # # 0 (# #)
      10) Save< 217 # 53 0 (0.6159420 0.3840580) *
      11) Save>=217 # 8 1 (0.2162162 0.7837838) *
  3) Bank>=2.5 # 23 1 (0.1446541 0.8553459) *
```

Some numbers are missing (denoted by #), please fill in these numbers.

- Find out all the missing numbers.
- Produce the classification table from this output and compute the error rate.
- Consider the terminal node 4). Compute support, confidence and capture.

## 4 Building Classification Tree

### 4.1 Impurity Measure

There are three measures of the impurity of the C-Tree:

- $Entropy(t) = - \sum_{i=0}^{c-1} p(i|t) \log_2 p(i|t)$  with  $0 \log_2 0 = 0$
- $Gini(t) = 1 - \sum_{i=0}^{c-1} p(i|t)^2$
- $Classificationerror(t) = 1 - \max_i [p(i|t)]$

Since the objective of C-tree is to find the rule such that the terminal nodes are as “pure” as possible, so for every splitting, we would like to find a variable  $v$  such that the goodness of splitting variable  $v$  is the largest.

### 4.2 Algorithm

The main idea to build a classification tree is to find the attribute that minimize entropy at one step, then after we splitting the dataset into several sub datasets, we find the best attributes for each sub dataset. Then rerun this procedure recursively.

---

**Algorithm 1:** Classify

---

**Input:** Data  
**Output:** Classification tree  
Initialize tree  
Find the best attribute that minimize the entropy.  
**foreach** *Sub dataset* **do**  
    subtree=Classify(sub dataset)  
    Append subtree on tree.  
**end**  
Return tree

---